

REPORT 4

Comparison and Interpretation

CNN – Medical Image Classification

Name: Selin Dyortkardeshler

Submission Date: May 14 2026

1. Introduction

This report constitutes the final analytical chapter of a four-part series documenting the development of a deep learning pipeline for binary classification of chest X-ray images into NORMAL and PNEUMONIA categories. The preceding reports covered data preparation and augmentation (Report 1), custom CNN design and iterative improvement from Variant A to Variant B (Report 2), and ResNet50 transfer learning with two-phase training (Report 3).

The objective of this report is to formally compare both modeling approaches, the custom CNN (Variant B) and the ResNet50 transfer learning model, across all relevant performance dimensions, interpret prediction errors, analyze the trade-offs involved, and deliver a final project conclusion with methodological reflection.

All evaluations are conducted on the same held-out test set of 624 images (234 NORMAL, 390 PNEUMONIA) to ensure fair, unbiased comparison.

2. Summary Comparison Table

Table 1 provides a comprehensive side-by-side comparison of all models developed throughout this project: Variant A (baseline CNN), Variant B (improved CNN), and the ResNet50 transfer learning model. Metrics are reported at each model's optimal threshold to ensure the fairest comparison of discriminative capacity.

Metric	Variant A(t=0.944)	Variant B(t=0.791)	ResNet50 TL(t=0.5)	Best Model
Accuracy	73.08%	84.62%	92.15%	ResNet50 ✓
Weighted F1	73.48%	85.43%	93.82%	ResNet50 ✓
PNEUMONIA: Recall	66.41%	83.59%	95.38%	ResNet50 ✓
NORMAL Recall	84.19%	88.03%	86.75%	Variant B
Precision	~75%	~85%	92.31%	ResNet50 ✓
ROC-AUC	0.784	0.925	0.974	ResNet50 ✓
Test Loss	0.801	0.382	0.283	ResNet50 ✓
Optimal Threshold	0.944	0.791	0.971*	Variant B†
Trainable Parameters	110,785	110,785	~23M	Variant B
Training Epochs	18	18	40 (2×20)	Variant B

* ResNet50's optimal threshold of 0.971 indicates probability calibration issues; the default threshold t=0.5 is used for primary comparison.

† Variant B's threshold of 0.791 is more clinically interpretable, despite lower overall performance.

Table 1. Full model comparison across all performance metrics.

3. Which Model Performs Best and Why

The ResNet50 transfer learning model is the clear overall winner, outperforming Variant B on every primary metric: +7.5 percentage points in accuracy, +8.4 points in F1, +11.8 points in pneumonia recall, and +0.049 in ROC-AUC. The reasons for this superiority are both architectural and methodological.

3.1 Why ResNet50 Outperforms the Custom CNN

ResNet50 was pre-trained on 1.2 million ImageNet images, giving it deeply learned feature hierarchies, edge detectors, texture classifiers, and shape recognizers before any medical image was shown to it. These representations transfer surprisingly well to chest X-rays because the low-level visual features (local contrast patterns, boundary gradients) that characterize consolidation and infiltrates in pneumonia are structurally similar to natural image features.

The custom CNN, by contrast, had to learn all features from only 5,216 training images a fundamentally harder optimization problem. Despite regularization improvements in Variant B, the network's 110,785 parameters could not achieve the representational richness of ResNet50's 23 million parameters, shaped by large-scale pre-training.

Additionally, the two-phase training strategy (frozen feature extraction followed by selective unfreezing) provided a stable optimization path that avoided destructive interference between the pre-trained weights and the new classification objective.

3.2 The CNN Advantage: Efficiency and Calibration

Despite lower absolute performance, Variant B holds meaningful advantages. Its 110,785 parameters are 200 times fewer than ResNet50, making it far more suitable for resource-constrained deployment environments such as mobile diagnostic tools or edge devices in low-resource healthcare settings.

Furthermore, Variant B's optimal threshold of 0.791 is considerably more interpretable than ResNet50's 0.971. A threshold near 0.5 indicates well-calibrated probabilities the model's confidence scores reflect genuine uncertainty. ResNet50's threshold of 0.971 suggests that its sigmoid outputs cluster near 1.0 regardless of true confidence, which complicates clinical decision-making and threshold interpretation.

4. Prediction Error Analysis

In binary classification, prediction errors fall into two categories: false negatives (FN) pneumonia cases classified as normal and false positives (FP) normal cases classified as pneumonia. The clinical and operational consequences of these two error types differ substantially.

4.1 Confusion Matrix — Custom CNN (Variant B, $t=0.791$)

The confusion matrix in Figure 1 (from Report 2) reflects Variant B's performance at its optimal threshold. The model produces 64 false negatives and 28 false positives. The ratio FN:FP = 64:28 $\approx$ 2.3:1, indicating that the model is approximately twice as likely to miss a pneumonia case as it is to misclassify a healthy patient.

4.2 Confusion Matrix — ResNet50 Transfer Learning (t=0.5)

Figure 1 and Figure 2 show ResNet50 confusion matrices at the default and optimal thresholds, respectively. The model achieves only 18 false negatives and 31 false positives, a dramatically lower false negative rate compared to Variant B. At the optimal threshold, (t=0.971), false positives drop to just 9, but false negatives rise sharply to 51, exceeding Variant B's count despite the overall superiority of the model.




Figure 1. ResNet50 confusion matrix at default threshold (t=0.5): 18 FN, 31 FP.




Figure 2. ResNet50 confusion matrix at optimal threshold (t=0.971): 51 FN, 9 FP.

4.3 Error Distribution Summary

Table 2 compares error counts across all three models at their respective reported thresholds.

Model / Threshold	True Positive	True Negative	False Negative	False Positive
Variant A (t=0.944)	259	197	131	37
Variant B (t=0.791)	326	206	64	28
ResNet50 (t=0.5)	372	203	18	31
ResNet50 (t=0.971)	339	225	51	9

Table 2. Prediction error counts across all models.

The most clinically significant observation is the progressive reduction in false negatives: from 131 (Variant A) to 64 (Variant B) to 18 (ResNet50 at t=0.5). This represents an 86% reduction in missed pneumonia cases across the project's development arc.

5. Error Distribution Analysis

5.1 Class-Level Error Rates

The error rates reveal a systematic pattern across all three models: they all struggle more with the NORMAL class than the PNEUMONIA class at default threshold settings. This is a direct consequence of training on a class-imbalanced dataset where 74.3% of training images are PNEUMONIA. Despite balanced class weighting, the models' decision boundaries are shifted towards predicting PNEUMONIA, producing lower NORMAL recall in Variant A (14.96%) and Variant B (71.79%) at t = 0.5.

5.2 Threshold-Dependent Error Behaviour

Figure 3 illustrates how the decision threshold controls the trade-off between PNEUMONIA recall and NORMAL recall for ResNet50. Two key observations emerge:

- At low thresholds ($t < 0.3$), PNEUMONIA recall approaches 100%, but NORMAL recall falls below 70%, meaning nearly all patients are flagged as pneumonia.
- At high thresholds ($t > 0.8$), NORMAL recall exceeds 90%, but PNEUMONIA recall declines, increasing the risk of missed diagnoses.
- The weighted F1 score peaks near $t=0.5$, confirming that the default threshold represents the best aggregate balance point for ResNet50.




Figure 3. Effect of decision threshold on ResNet50 performance metrics. The default threshold ($t=0.5$) maximises weighted F1; the optimal threshold ($t=0.971$) maximises Normal Recall.

5.3 ROC Curve Analysis

Figure 4 shows the ResNet50 ROC curve with an AUC of 0.974. The curve rises steeply toward the top-left corner, indicating that the model correctly ranks pneumonia cases above normal cases with very high probability. An AUC of 0.974 means that for any randomly selected (pneumonia or normal) pair, the model assigns a higher score to the pneumonia case 97.4% of the time a strong discriminative result.

By comparison, Variant B achieved an AUC of 0.925 and Variant A only 0.784. The 0.049 AUC gain from Variant B to ResNet50 represents a meaningful improvement in the model's underlying ranking ability, independent of threshold selection.




Figure 4. ROC Curve — ResNet50 Transfer Learning (AUC = 0.974). Red dot marks the optimal threshold ($t=0.971$).

6. Performance on Challenging Cases

In classification tasks, 'peak demand' scenarios correspond to the most difficult or ambiguous cases, those that stress the model's discriminative capacity most severely. In the context of chest X-ray classification, these correspond to three identifiable subgroups:

6.1 Ambiguous Cases: Near-Threshold Predictions

Cases where the model assigns probabilities near the decision boundary (0.4–0.6) $t=0.5$ represent the highest uncertainty region. For ResNet50, the sigmoid output distribution clusters near 0 and 1, with relatively few predictions in the mid-range. This is a sign of high model confidence but poor calibration; the model is extremely decisive even on cases where radiological ambiguity is expected.

Variant B, with a more graduated probability distribution (optimal threshold 0.791), handles near-boundary cases with greater interpretable uncertainty, which may actually be more appropriate in clinical consultation settings where a radiologist needs to know the model's confidence level.

6.2 Class-Imbalance Stress: Minority Class Performance

The NORMAL class constitutes only 37.5% of the test set (234 images) and 25.7% of the training set. Performance on the NORMAL class is, therefore, the most demanding test of the models' resistance to majority-class bias. Across all models at default threshold:

Model	NORMAL Recall	PNEUMONIA: Recall	Gap
Variant A (t=0.5)	14.96%	92.82%	77.86%
Variant B (t=0.5)	71.79%	92.31%	20.52%
ResNet50 (t=0.5)	86.75%	95.38%	8.63%

Table 3. Recall gap between PNEUMONIA and NORMAL classes at default threshold.

The recall gap narrows progressively from 77.86% (Variant A) to 8.63% (ResNet50), demonstrating that transfer learning substantially resolves the class-imbalance challenge that plagued the custom CNN approaches.

6.3 Training Stability Under Extended Training

Phase 2 training curves (Figure 5) reveal performance under extended optimization analogous to a stress period in demand forecasting. After epoch 12, training accuracy continues to improve (~99%) while validation accuracy plateaus and begins to oscillate (~97%). This divergence is a mild overfitting signal, indicating that the model would have benefited from early stopping at epoch 12, which would have yielded a test loss approximately 28% lower than the final checkpoint.




Figure 5. Phase 2 training curves. Divergence between train and validation loss after epoch 12 indicates mild overfitting.

7. Trade-Off Discussion

7.1 Accuracy vs Computational Cost

The performance improvement from Variant B to ResNet50 comes at a substantial computational price. ResNet50 has approximately 200 times more parameters (23M vs 110K), requires significantly more memory, and takes roughly 4× longer to train per epoch on consumer hardware. Table 4 summarizes this trade-off:

Dimension	Variant B (CNN)	ResNet50 (TL)
Parameters	110,785	~23,000,000
Accuracy	84.62%	92.15%
ROC-AUC	0.925	0.974
Training time	Faster	~4× slower
Memory footprint	Low	High
Deployment cost	Low	High
Calibration	Better (t=0.791)	Poor (t=0.971)

Table 4. Accuracy vs computational cost trade-off summary.

For high-resource hospital settings, where diagnostic accuracy is paramount and computational resources are abundant, ResNet50 is the clear choice. For mobile or low-bandwidth deployment such as rural clinics in developing regions where pneumonia is particularly prevalent, Variant B's efficiency may outweigh its accuracy disadvantage.




Figure 6. Custom CNN (Variant B) vs. ResNet50 Transfer Learning — all four primary metrics. ResNet50 leads on all measures.

7.2 Simplicity vs Temporal (Iterative) Modelling

This project adopted an iterative modeling strategy: each model version explicitly addressed identified failures from the previous one. Variant A revealed the critical importance of validation set quality; Variant B demonstrated that architectural simplicity with proper regularisation

achieves competitive results; ResNet50 established the ceiling achievable with pre-trained representations.

The simplicity of the custom CNN provides important interpretive advantages: the model is fully transparent, its parameters can be inspected, and its decision behavior is more predictable. ResNet50, while more powerful, operates as a near-black box its 23 million parameters do not lend themselves to straightforward inspection or debugging.

In a regulated clinical environment, model transparency is not merely a convenience but a regulatory requirement. Explainability tools such as Grad-CAM can partially address this limitation for ResNet50, but they add complexity to the deployment pipeline that Variant B avoids entirely.

7.3 Sensitivity vs Specificity Clinical Trade-off

In medical screening, the fundamental trade-off is between sensitivity (recall) and specificity. Missing a pneumonia case (false negative) can delay treatment and worsen patient outcomes. Falsely flagging a healthy patient (false positive) leads to unnecessary further examination, patient anxiety, and healthcare costs.

At default threshold:

- ResNet50: Pneumonia Recall = 95.38%, Normal Recall = 86.75% strong screening performance.
- Variant B: Pneumonia Recall = 92.31%, Normal Recall = 71.79%, higher false alarm rate.

For a primary screening application, ResNet50 at=0.5 is recommended, as it achieves the best balance. If a secondary confirmatory use case is required (e.g., reducing unnecessary hospital admissions), ResNet50 t=0.971 reduces false positives to 9 but should only be deployed with physician oversight given the increased false negative count.

8. Discussion

8.1 Methodological Contributions

This project's most significant methodological contribution is the explicit identification and correction of a critical data preprocessing error: the Kaggle dataset's original 16-image validation set rendered callbacks non-functional and produced a model (Variant A) with a 30-point train-test accuracy gap. Recognizing this failure and correcting it via a stratified 20% validation split produced a 21-point accuracy improvement in Variant B without any change to model architecture, a striking demonstration that data pipeline quality can dominate model design quality.

The two-phase transfer learning strategy also represents a sound methodological choice. By freezing the backbone in Phase 1, the classification head was allowed to stabilize before the computationally expensive fine-tuning phase, reducing the risk of catastrophic forgetting of the pre-trained representations.

8.2 Comparison of Approaches

Three distinct modeling philosophies were explored:

- Tabula rasa learning (Variant A/B): Training from random initialization on domain-specific data. Advantages: full transparency, low resource cost. Disadvantages: data-hungry, prone to overfitting on small datasets.
- Feature extraction transfer learning (Phase 1): Frozen backbone, trainable head. Advantages: rapid convergence, leverages ImageNet representations. Disadvantages: upper-bounded by head capacity.
- Fine-tuning transfer learning (Phase 2): Selective unfreezing. Advantages: domain adaptation of deep features. Disadvantages: risk of overfitting if learning rate is too high; requires careful epoch management.

The progression from Variant A (AUC=0.784) to Variant B (AUC=0.925) to ResNet50 (AUC = 0.974) clearly illustrates the incremental value of each methodological improvement.

8.3 Limitations

Several limitations must be acknowledged:

- Dataset scope: The Chest X-Ray Pneumonia dataset consists exclusively of pediatric patients aged 1–5 from a single Chinese medical center. Generalisability to adult patients, different ethnic groups, or radiological equipment from other institutions remains unvalidated.
- Probability calibration: Both ResNet50 ($t=0.971$) and Variant A ($t=0.944$) exhibit poor probability calibration. Without calibration post-processing (e.g., Platt scaling or temperature scaling), the model's confidence scores should not be interpreted as true probabilities. This is a critical limitation for any clinical deployment.
- Validation set size: Despite the correction in Variant B, the 20% stratified split produces only ~1,043 validation images drawn from the same distribution as training data. This may produce overly optimistic validation metrics. K-fold cross-validation would provide more robust estimates.
- Overfitting in Phase 2: Training beyond epoch 12 in ResNet50's fine-tuning phase introduced mild overfitting. Early stopping was not applied in this project; its inclusion would have produced a marginally better test checkpoint.
- No external validation: The test set is drawn from the same source as the training data. True clinical validation requires evaluation on an independent dataset from a different institution and imaging protocol.

9. Final Conclusions

This project developed and evaluated a complete deep learning pipeline for pneumonia detection from chest X-ray images, progressing through three modeling stages and four reports. The key findings are summarized below.

9.1 Primary Finding

The ResNet50 transfer learning model is the best-performing system, achieving an ROC-AUC of 0.974, an accuracy of 92.15%, and a pneumonia recall of 95.38% at the default threshold. It outperforms the custom CNN (Variant B) by +5.8% accuracy and +4.9% AUC. This confirms the core hypothesis of the project: that pre-trained representations from large-scale natural image

datasets transfer effectively to the binary classification of chest X-rays, even across domain boundaries.

9.2 Development Arc

The iterative development produced consistent, measurable improvements at each stage:

- Variant A → Variant B: +21% accuracy through validation pipeline correction and regularization.
- Variant B → ResNet50: +7.5% accuracy through transfer learning and two-phase fine-tuning.
- Overall: +28.5% accuracy improvement from naive baseline (Variant A, 63.62%) to final model (ResNet50, 92.15%).
- False negative reduction: 131 (Variant A) → 64 (Variant B) → 18 (ResNet50), an 86% improvement.

9.3 Clinical Deployment Recommendation

For clinical deployment, the following strategy is recommended based on project findings:

- Primary screening: ResNet50 at $t=0.5$ maximizes pneumonia recall (95.38%) to minimize missed cases.
- Confirmatory filtering: ResNet50 at $t=0.971$ maximises precision (97.41%) to reduce unnecessary hospitalisations, but only under physician oversight given increased false negative rate.
- Resource-constrained settings: Variant B at $t=0.5$ or $t=0.791$ is lightweight and deployable on edge devices, with acceptable recall of 92.31%.
- Required before deployment: Probability calibration (Platt scaling or temperature scaling), Grad-CAM explainability integration, and validation on an independent external dataset.

9.4 Overall Assessment

This project demonstrates that robust pneumonia classification is achievable with deep learning on publicly available data and that methodological rigor, particularly in data pipeline design and validation strategy, is at least as important as model architecture choice. The transfer learning approach achieves near-radiologist-level discriminative performance (AUC=0.974) on this dataset, representing a meaningful contribution to the broader goal of AI-assisted medical diagnosis.

Project Summary — Key Results

Best Model: ResNet50 Transfer Learning | ROC-AUC: 0.974 | Accuracy: 92.15%
Pneumonia Recall: 95.38% | F1 Score: 93.82% | False Negatives (test): 18 / 390
Improvement over naive baseline: +28.5% accuracy | FN reduction: 86%

References

- [1] Mooney, P. (2018). Chest X-Ray Images (Pneumonia). Kaggle. <https://www.kaggle.com/datasets/paultimothymooney/chest-xray-pneumonia>
- [2] Kermany, D. S., et al. (2018). Identifying Medical Diagnoses and Treatable Diseases through Image-Based Deep Learning. *Cell*, 172(5), 1122–1131.
- [3] He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep Residual Learning for Image Recognition. *CVPR* 2016.
- [4] Chollet, F. (2021). *Deep Learning with Python* (2nd ed.). Manning Publications.
- [5] Srivastava, N., et al. (2014). Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *JMLR*, 15(1), 1929–1958.
- [6] Kingma, D. P., & Ba, J. (2015). Adam: A Method for Stochastic Optimization. *ICLR* 2015.
- [7] Platt, J. (1999). Probabilistic Outputs for Support Vector Machines. *Advances in Large Margin Classifiers*.
- [8] Selvaraju, R. R., et al. (2017). Grad-CAM: Visual Explanations from Deep Networks. *ICCV* 2017.